ASSIGNMENT - 4 : DATA WAREHOUSING LAB

Submitted by : Shivie Saksenaa

1. Objective

The objective of this lab was to understand how **Google Cloud Storage (GCS)** and **BigQuery** integrate to provide a complete cloud-based data storage and analytics solution.
The goal was to:

- Store data in **GCS**,
- Query it directly using an **external table** in **BigQuery**, and
- Compare performance with a **managed (partitioned) BigQuery table**.

2. Tools and Services Used

Tool	Purpose
Google Cloud Storage (GCS)	Store raw CSV data securely
BigQuery	Data warehousing and SQL-based analytics
GCP Project ID	clean-equinox-472523-c4
Dataset	ecommerce_dw
Region	us-east1
Dataset used	Ecommerce_Orders.csv

3. Steps Performed

Step 1: Upload data to Google Cloud Storage

1. Created bucket : shivie-lab4-data-bucket
2. Added folder : raw/
3. Uploaded file : ecommerce_orders.csv
4. Final GCS path: gs://shivie-lab4-data-bucket/raw/ecommerce_orders.csv

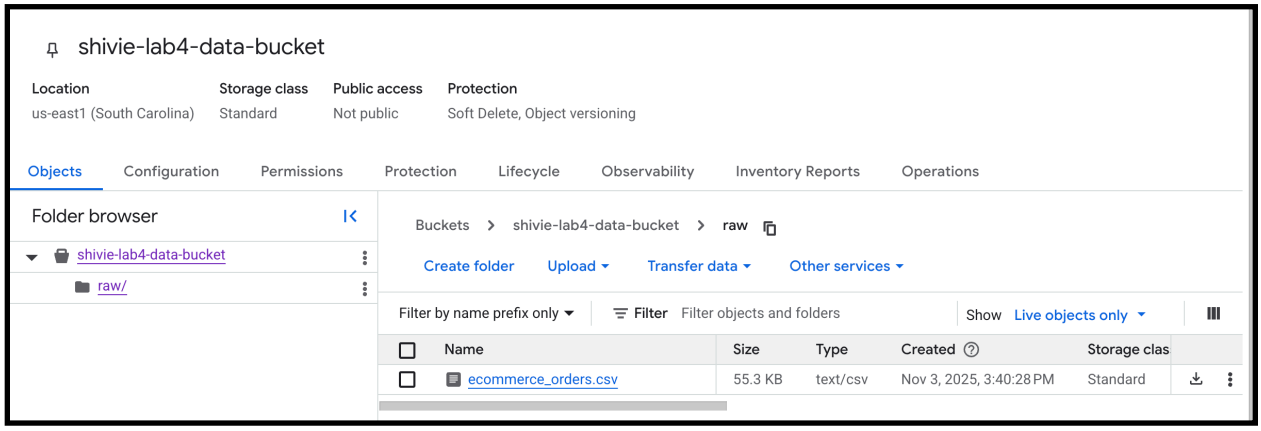

Step 2: Created External Table in BigQuery

1. In BigQuery Console , I created the Table **orders.ext**
    - **Source:** Google Cloud Storage
    - **URI:** gs://shivie-lab4-data-bucket/raw/ecommerce_orders.csv
    - **File Format:** CSV
    - **Auto Detect Schema:** Enabled
    - **Destination Dataset:** ecommerce_dw
    - **Table Type:** External Table

Result : Table was created : clean-equinox-472523-c4.ecommerce_dw.orders_ext

Step 3: Created Managed (Partitioned) Table

1. In BigQuery, I created the Table **orders_partitioned**
    - **Source:** Same GCS file
    - **Destination Table:** ecommerce_dw.orders
    - **Table Type:** Native (Managed)
    - **Partition By:** order_date
    - **Cluster By:** state, category

Result: Table Created: clean-equinox-472523-c4.ecommerce_dw.orders

AIM: Understand the difference in performance of External Table and Native table.
Query Tested: on both the tables - External and the Native One

```
SELECT
  DATE_TRUNC(order_date, MONTH) AS month,
  ROUND(SUM(line_amount_usd), 2) AS total_revenue
FROM `clean-equinox-472523-c4.ecommerce_dw.orders_ext`
GROUP BY month
ORDER BY month;
```

4. Result

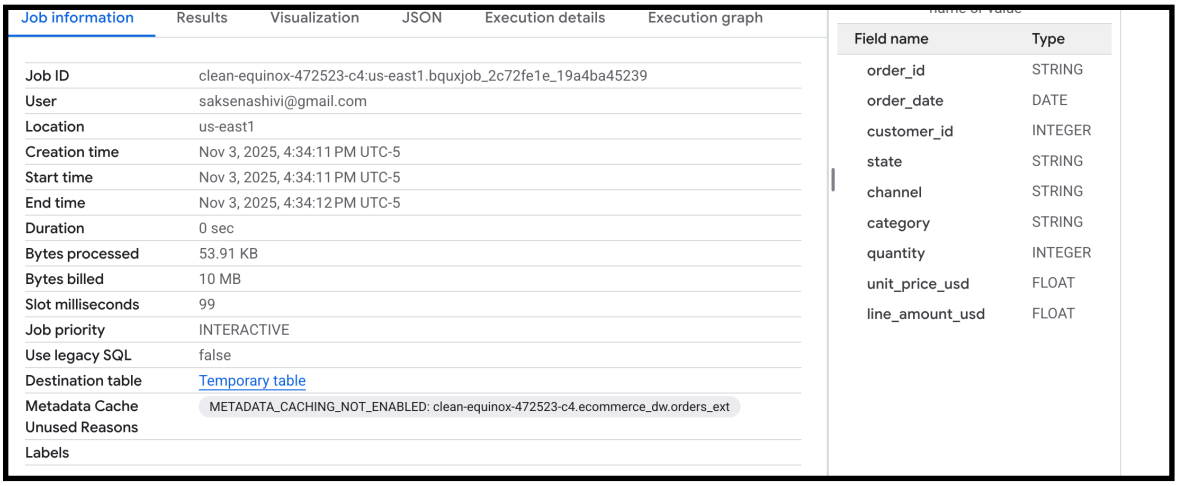

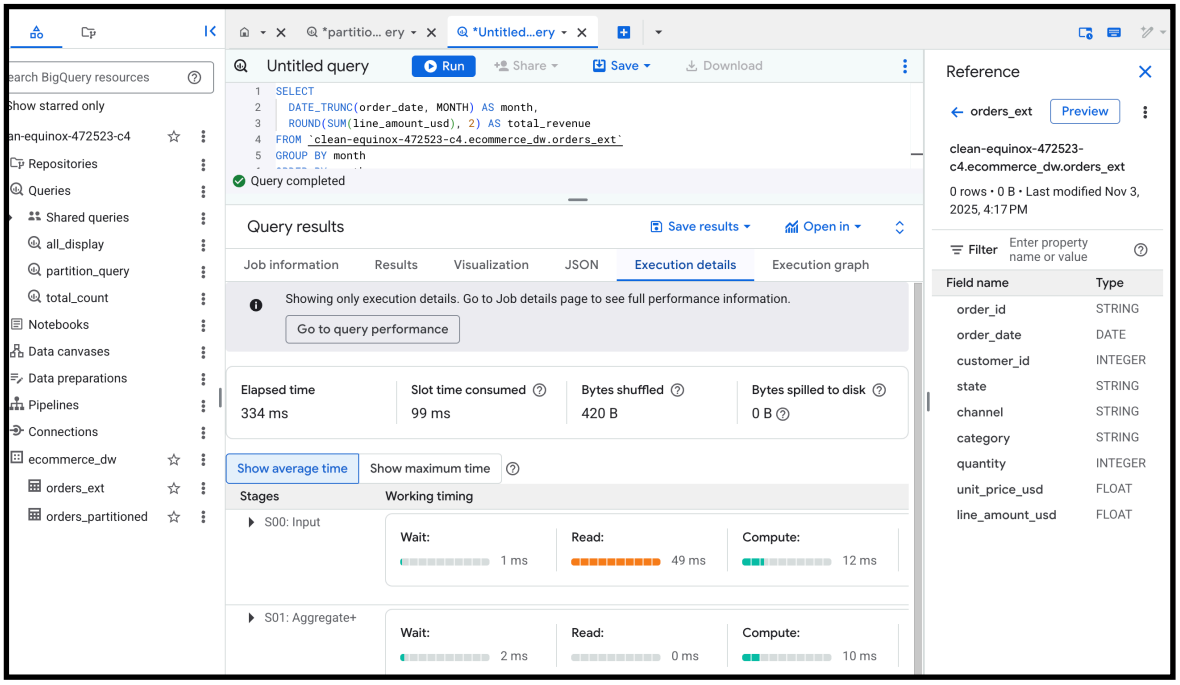

Observation and Inference – External Table (orders_ext)

- When I ran the query on the external table, BigQuery read the data directly from my GCS file (gs://shivie-lab4-data-bucket/raw/ecommerce_orders.csv).
  The job processed 53.91 KB and took around 334 ms. Since the table only references the CSV, BigQuery had to scan the entire file every time I ran the query.
  Caching, partitioning, and clustering weren't available, and most of the time was spent reading data from GCS.
- From this, I inferred that external tables are great for quick, ad-hoc analysis of raw files but not ideal for repeated or large analytical queries. They provide flexibility without needing to load data, but performance is slower compared to a managed table

Managed Table

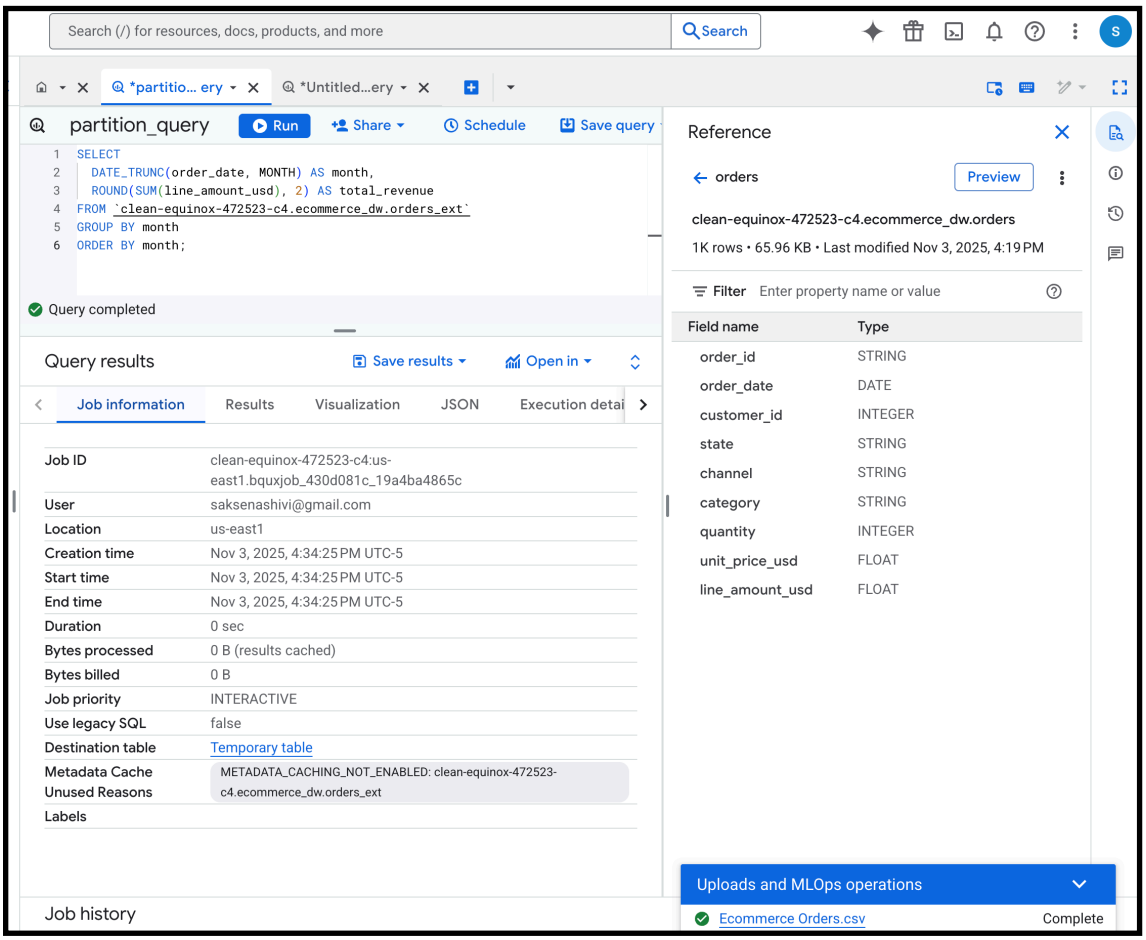

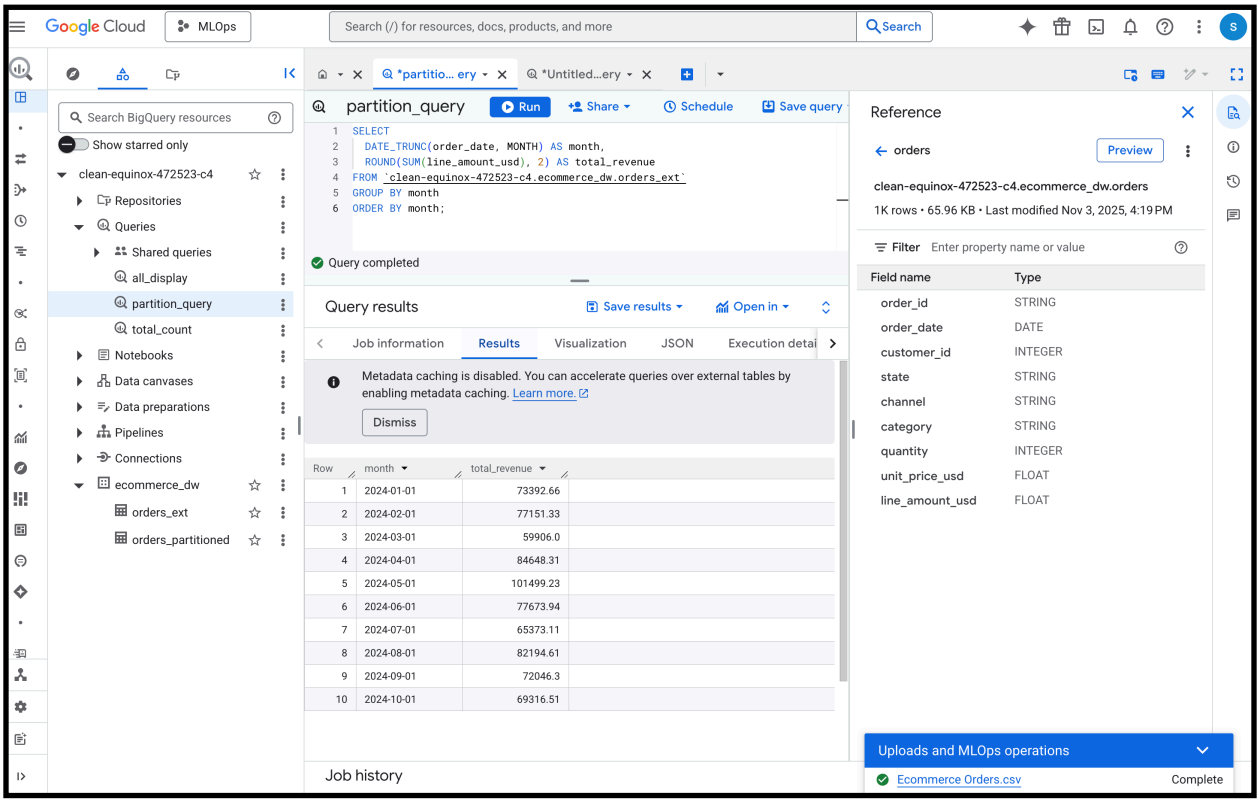

Observation and Inference – Managed Table (orders)

- When I ran the same query on the managed, partitioned table, BigQuery used its internal storage instead of reading from GCS.
  The job showed 0 bytes processed and finished instantly because the data was already stored in BigQuery and cached.
  Since this table is partitioned by order_date and clustered by state and category, BigQuery only scans the required partitions, making queries faster and more efficient.

- From this, I inferred that managed tables are much better for repeated and large-scale analytics. They support partitioning, clustering, and caching, which reduce scan cost and improve performance compared to external tables.